

# Transport Mode Detection - Nimish Agarwal (170440)

## Setup

In [1]:

```
# Mounting Drive
from google.colab import drive
drive.mount('/content/drive')
%cd /content/drive
```

Mounted at /content/drive  
/content/drive

In [2]:

```
# Importing Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime, timedelta
import time
import io
```

In [3]:

```
MAIN_DIR = 'My Drive/4th year Stuff/CE784A/TMD_Dataset/'
DATASET_DIR = MAIN_DIR + 'cleaned.csv'
```

## Part 1 - Reading and Exploring Dataset

In [4]:

```
df = pd.read_csv(DATASET_DIR)
df.head()
```

Out[4]:

	user	timestamp	x	y	z	class
0	a2d80ed662f34d32951eb1c6ed076c313e358b73	2018-06-04 16:26:55.053	0.78	-9.13	-3.74	bus
1	a2d80ed662f34d32951eb1c6ed076c313e358b73	2018-06-04 16:26:55.111	0.79	-9.11	-3.75	bus
2	a2d80ed662f34d32951eb1c6ed076c313e358b73	2018-06-04 16:26:55.169	0.80	-9.12	-3.75	bus
3	a2d80ed662f34d32951eb1c6ed076c313e358b73	2018-06-04 16:26:55.228	0.78	-9.14	-3.76	bus
4	a2d80ed662f34d32951eb1c6ed076c313e358b73	2018-06-04 16:26:55.286	0.83	-9.12	-3.80	bus

In [5]:

```
print(df.info())
print('classes: ', df['class'].unique())
# Unique users
print('Unique users: ', len(df['user'].unique()))
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5653053 entries, 0 to 5653052
Data columns (total 6 columns):
#   Column      Dtype
---  -
0    user      object
1    timestamp  object
2    x          float64
```

```

3     y          float64
4     z          float64
5     class      object
dtypes: float64(3), object(3)
memory usage: 258.8+ MB
None
classes: ['bus' 'walk' 'car' 'bike' 'train' 'e-bike']
Unique users: 32

```

In [6]:

```
df.describe()
```

Out[6]:

	x	y	z
count	5.653053e+06	5.653053e+06	5.653053e+06
mean	1.499442e+00	1.483885e+00	2.484874e+00
std	4.657316e+00	6.262899e+00	5.800348e+00
min	-7.321000e+01	-7.840000e+01	-7.844000e+01
25%	-1.300000e+00	-1.790000e+00	-9.600000e-01
50%	7.100000e-01	2.130000e+00	3.500000e+00
75%	4.650000e+00	6.260000e+00	7.320000e+00
max	7.840000e+01	7.834000e+01	7.840000e+01

## Part 2 - Number of Unique Sequences

In [7]:

```
df['timestamp'] = pd.to_datetime(df['timestamp'])
df_temp = df[['user', 'class', 'timestamp']]
```

In [8]:

```

start=time.time()
# get time difference by Row[i] - Row[i-1]
temp = (abs(df_temp['timestamp'][1:].reset_index(drop=True) - df_temp['timestamp'][:-1])
> timedelta(seconds = 10))
# Changing mode
temp = temp | (df_temp['class'][1:].reset_index(drop=True) != df_temp['class'][:-1])
df_temp['gap'] = pd.concat([pd.DataFrame(data = [True]), temp ]).reset_index(drop=True)
)

```

/usr/local/lib/python3.6/dist-packages/ipykernel\_launcher.py:6: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

### Per user per trasportation mode, No. of Unique Sequences

In [9]:

```

# Grouping dataset based on
df_temp2 = (df_temp.groupby(['user', 'class', 'gap']).count()).reset_index()
# Findng no. of unique sequences
df_temp2 = df_temp2[df_temp2.gap != False]
# taking 3 columns for simplicity
df_temp2 = df_temp2[['user', 'class', 'timestamp']]
df_temp2 = df_temp2.rename(columns={"timestamp": "Sequences"})
df_temp2 = df_temp2.reset_index(drop=True).groupby(["user", "class", "Sequences"]).count()

```

```
#dropping index
pd.set_option('display.max_rows', len(df_temp2))
display(df_temp2)
pd.reset_option('display.max_rows')
end=time.time()
print("Time to run the code :",end-start)
```

	user	class	Sequences
	a2d80ed662f34d32951eb1c6ed076c313e358b73	bus	13
	a526f3566e9c9024dfa7378eb4291d787a09fd37	car	15
		walk	19
	a59868c6eb3645eedbb343ce8c336ec6f2ef2324	bike	39
		bus	27
		car	68
		walk	12
	a92dee88f61123f923dcceec01eeecf1a81953b36	bus	4
	ac4c17afeb69b39169eb301ab592696a8f353976	car	20
		walk	58
	adaaae1a67ea9e43abd60ba945eccda0cb8821e0	bus	22
		car	17
		walk	10
	b138d165100ef60bc793cac143742eb5aea4d6ba	car	9
	b45157069942d01310c3e7b74034166717bb25f9	car	3
		walk	2
	b7b165e5637b5a0226068d907748f4bbfc61a320	car	194
		walk	10
	c453226e3616ae821cdcb38f38481c2a20f2482f	bike	169
		bus	291
	c5702d34b238fe68683f818e82cd3a3cd8a16366	bike	15
		bus	14
		walk	129
	ca7950f223a8037b897d0547075dc138f9e43b20	walk	3
	cace4ec0999436917986b4fa6e9317262c897bc2	car	72
	cbde60baea002b694ecf2a3ff2d95be16b00efe1	bus	4
	ce39f5d0705695fcd70a04ba6d84ac6beecd6f9c	bus	22
		car	80
	d429974540bfd38c3367fe9f0c8682775ff4fa18	bus	6
		car	132
		walk	2
	d7a1230d94f91a32cc079809748e52e8a4a6a22f	bike	152
		train	13
		walk	6
	d7dd12d83c81574137f858034b99f4cc83ab0718	car	147
	d8c047eaaee204b7b5cd71e2d67308b87b038ed3	car	12
		walk	24
	dc0bdce306ec3b624fe0e6ecd1ffbd82cb970120	bike	115
		car	32

	user	walk class	Sequences	5
	dd82e3df4bebc74ed6b67877be79e29f401c16a3	car		68
		walk		2
	dde95e125d89843f7032baa734ee4d34ec775aaf	bus		5
		car		2
		walk		8
	de9892b879c83ea3d24fb4560873107cc4e86d48	car		114
		walk		36
	dfcfc0404691b73b69884073159f90843f2ac35b	bus		50
		car		108
		walk		155
	e429a95c532f1117130c11e4a18379d84fa4ffa9	bus		40
		car		39
	eb9e7854290fd6ea9ebaf448b640fc1f1dbeb076	bus		2
		train		1
		walk		2
	ecfb0929250fb6dda66a4065441230ab27f094e5	car		158
		e-bike		16
		train		1
	ed623d28c1e0071632a6110b8f8ed93f8af78b99	bus		10
		car		117
		walk		4
	f1b7331b66e404c11eebb22933e733117bbb12c9	bike		172
		car		139
		walk		73
	f5edd999397145a2ec1b244226fc83f99631760c	bus		16
		walk		13
	f7ae1ce141c26db40ea8b090fb568a0c965310aa	car		2
	faae5be800be2dfa897eea0bd2e5988cd53c4ec0	bike		136
		car		10
		walk		35

Time to run the code : 1.8916699886322021

## Part 3 - Time Window Partition

In [10]:

```
df = pd.read_csv(DATASET_DIR)
df['timestamp'] = df['timestamp'].astype('datetime64[s]')
```

In [11]:

```
start=time.time()
def helper_func(temp):
    temp = temp["timestamp"]
    temp = pd.DataFrame(data = temp).set_index("timestamp")
    temp["val"] = np.ones(len(temp))
    temp = temp.groupby(temp.index).count()
    temp = temp.resample('5s', origin='start').count()
    return len(temp[temp.val != 0])
```

In [14]:

```
# preprocessing data
df_temp = df[['user', 'class', 'timestamp']]
temp = (abs(df_temp['timestamp'][1:].reset_index(drop=True) - df_temp['timestamp'][:-1]))
> timedelta(seconds = 10))
df_temp['gap'] = pd.concat([pd.DataFrame(data = [True]), temp]).reset_index(drop=True)
df_temp2 = df_temp[df_temp["gap"] == True]
df_temp2 = (df_temp2.groupby(['user', 'class', 'gap']).apply(lambda x : x)).reset_index().drop(['gap'], axis=1)
df_temp2["end_timestamp"] = df_temp.reset_index()["timestamp"].iloc[(df_temp2["index"]-1).to_numpy()[1:]].reset_index(drop=True)
df_temp2["end_timestamp"].iloc[-1] = df["timestamp"].iloc[-1]
df_temp2[">than5sec"] = ((df_temp2["end_timestamp"] - df_temp2["timestamp"]) >= timedelta(seconds = 5))
df_temp2["start_index"] = (df_temp2["index"].to_numpy())
df_temp2["end_index"] = np.concatenate((((df_temp2["index"]-1).to_numpy()[1:]), [len(df)-1]), axis = 0)
# starting and ending of each sequence
df_temp3 = df_temp2[df_temp2[">than5sec"] == True]
a = list(range(df_temp3["start_index"].iloc[0], df_temp3["end_index"].iloc[0]-df_temp3["start_index"].iloc[0]+1))
for i in range(1, len(df_temp3)):
    a.extend(list(range(df_temp3["start_index"].iloc[i], df_temp3["end_index"].iloc[i]+1)))
df6 = df.iloc[a]

df_temp3 = pd.DataFrame(df6.groupby(["user", "class"]).apply(helper_func))
df_temp3 = df_temp3.reset_index().rename(columns={0: "#5_SecWindowSequences"}).groupby(["user", "class", "#5_SecWindowSequences"]).count()
df_temp3 = df_temp3.reset_index().groupby(["class"]).sum().reset_index()
display(df_temp3)
end=time.time()
print("Time to run the code :", end-start)
```

/usr/local/lib/python3.6/dist-packages/ipykernel\_launcher.py:4: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

after removing the cwd from sys.path.

/usr/local/lib/python3.6/dist-packages/pandas/core/indexing.py:670: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

iloc.\_setitem\_with\_indexer(indexer, value)

	class	#5_SecWindowSequences
0	bike	2958
1	bus	10543
2	car	20419
3	e-bike	78
4	train	594
5	walk	5825

Time to run the code : 38.75357675552368

## Part 4 - Feature Engineering

### Extracting statistical features

In [15]:

```
def feature_extraction(temp):  
    return temp.set_index("timestamp").resample('5s', origin='start').agg({'x': ['mean', 'max', 'min', 'std'], 'y': ['mean', 'max', 'min', 'std'], 'z': ['mean', 'max', 'min', 'std']})
```

In [16]:



```
df_temp3 = pd.DataFrame(df6.groupby(["user", "class"]).apply(feature_extraction)).reset_index().dropna()  
df_temp3.columns = [''.join(col).strip() for col in df_temp3.columns.values]  
df_temp3.head()  
print(df_temp3.shape)
```



(40366, 15)





## Box Plots

In [34]:




```
# Box Plots  
fig, axes = plt.subplots(nrows=2, ncols=1, figsize=(20, 12*12))  
for i in range(3, 15):  
    sns.boxplot(x=df_temp3["class"], y=df_temp3.iloc[:, i], data=df_temp3, ax=axes[i-3])
```















## Part 5 - Creating Balanced Dataset

In [18]:

```
# Choosing 78 rows from each mode randomly
MinDataPoints = min(df_temp3.groupby(['class']).apply(lambda x: len(x)))
print("Minimum Data Points in any Transportatation Mode => ", MinDataPoints)

def fun2(temp):
    return temp.sample(n = 78)
df_new = df_temp3.groupby(["class"]).apply(fun2)
df_new = df_new.reset_index(drop=True).iloc[:,1:].drop('timestamp',axis = 1).rename(columns={'class': 'target'})
pd.DataFrame(df_new.groupby(["target"]).apply(lambda x: len(x))).rename(columns={0: "#Data Points"}).reset_index()
```

Minimum Data Points in any Transportatation Mode => 78

Out[18]:

	target	#DataPoints
0	bike	78
1	bus	78
2	car	78
3	e-bike	78
4	train	78
5	walk	78

In [19]:

```
df_new.shape
```

Out[19]:

```
(468, 13)
```

## part 6 - Splitting dataset

### Splitting

In [20]:

```
X = df_new.drop('target', axis=1)
Y = df_new['target']
```

In [21]:

```
from sklearn.model_selection import train_test_split

x, x_test, y, y_test = train_test_split(X, Y, test_size=0.2, train_size=0.8)
x_train, x_cv, y_train, y_cv = train_test_split(x, y, test_size = 0.25, train_size = 0.75)
```

In [22]:

```
print('No. of Datapoints:')
print('Training = ', len(x_train))
print('Cross Validation = ', len(x_cv))
print('Testing = ', len(x_test))
```

```
No. of Datapoints:
Training = 280
Cross Validation = 94
Testing = 94
```

## Part 7 - ML Model

### LogRegr., SVM, decision tree and random forest classifier

In [23]:

```
# Turn the values into an array for feeding the classification algorithms.
x = x.values
x_train = x_train.values
x_test = x_test.values
x_cv = x_cv.values
y = y.values
y_train = y_train.values
y_test = y_test.values
y_cv = y_cv.values
```

In [24]:

```

from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import precision_score, recall_score, f1_score, roc_auc_score, accuracy_score, classification_report, confusion_matrix
from sklearn.model_selection import cross_val_score, KFold
from collections import Counter
from sklearn.pipeline import make_pipeline
import collections
import warnings
warnings.filterwarnings("ignore")

```

In [25]:

```

classifiers = {
    "LogisiticRegression": LogisticRegression(),
    "Support Vector Classifier": SVC(kernel='linear', C=1),
    "Decision Tree Classifier": DecisionTreeClassifier(),
    "Random Forest Classifier": RandomForestClassifier()
}

classrep = []
confusionmatrices = {}
predicts = {}
kfold = KFold(n_splits=4, random_state=42)
for key, classifier in classifiers.items():
    classifier.fit(x_train, y_train)
    results = cross_val_score(classifier, x, y, cv = kfold)
    print("Classifier: "+classifier.__class__.__name__+" has a cross val score Accuracy:
%.3f%% (%.3f%%)" % (results.mean()*100.0, results.std()*100.0))
    predictions = classifier.predict(x_test)
    print('Classification Report of ',key,':\n',classification_report(y_test, predictions))
    predicts[key] = predictions
    classrep.append(classification_report(y_test, predictions));
    confusionmatrices[key] = confusion_matrix(y_test,predictions);

```

Classifier: LogisticRegression has a cross val score Accuracy: 65.763% (2.441%)

Classification Report of LogisiticRegression :

	precision	recall	f1-score	support
bike	0.47	0.39	0.42	18
bus	0.35	0.32	0.33	19
car	0.67	0.50	0.57	20
e-bike	0.71	0.94	0.81	18
train	0.67	1.00	0.80	8
walk	0.45	0.45	0.45	11
accuracy			0.56	94
macro avg	0.55	0.60	0.57	94
weighted avg	0.55	0.56	0.55	94

Classifier: SVC has a cross val score Accuracy: 68.706% (2.451%)

Classification Report of Support Vector Classifier :

	precision	recall	f1-score	support
bike	0.42	0.44	0.43	18
bus	0.47	0.47	0.47	19
car	0.62	0.50	0.56	20
e-bike	0.77	0.94	0.85	18
train	0.89	1.00	0.94	8
walk	0.33	0.27	0.30	11
accuracy			0.59	94
macro avg	0.59	0.61	0.59	94
weighted avg	0.57	0.59	0.57	94

Classifier: DecisionTreeClassifier has a cross val score Accuracy: 59.326% (6.660%)

Classification Report of Decision Tree Classifier :

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

bike	0.50	0.50	0.50	18
bus	0.47	0.42	0.44	19
car	0.57	0.60	0.59	20
e-bike	0.85	0.94	0.89	18
train	0.83	0.62	0.71	8
walk	0.58	0.64	0.61	11
accuracy			0.62	94
macro avg	0.63	0.62	0.62	94
weighted avg	0.61	0.62	0.61	94

Classifier: RandomForestClassifier has a cross val score Accuracy: 76.733% (2.917%)

Classification Report of Random Forest Classifier :

	precision	recall	f1-score	support
bike	0.56	0.56	0.56	18
bus	0.71	0.53	0.61	19
car	0.71	0.75	0.73	20
e-bike	0.85	0.94	0.89	18
train	0.89	1.00	0.94	8
walk	0.67	0.73	0.70	11
accuracy			0.72	94
macro avg	0.73	0.75	0.74	94
weighted avg	0.72	0.72	0.72	94

## Artificial Neural Network

In [26]:

```
# Implementing ANN
# multi-class classification with Keras
from keras.models import Sequential
from keras.layers import Dense
from keras.wrappers.scikit_learn import KerasClassifier
from keras.utils import np_utils
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
from sklearn.preprocessing import LabelEncoder
from sklearn.pipeline import Pipeline
```

In [27]:

```
X = df_new.drop('target', axis=1)
Y = df_new['target']
```

In [28]:

```
encoder = LabelEncoder()
encoder.fit(Y)
encoded_Y = encoder.transform(Y)
# convert integers to dummy variables (i.e. one hot encoded)
dummy_y = np_utils.to_categorical(encoded_Y)
```

In [29]:

```
from sklearn.model_selection import train_test_split

x, x_test, y, y_test = train_test_split (X, dummy_y, test_size=0.2, train_size=0.8)
```

In [30]:

```
print(x.shape)
print(y.shape)
```

```
(374, 12)
(374, 6)
```

In [31]:

```
n_inputs = x.shape[1]
y_out = y.shape[1]

# Defining Model
def baseline_model():
    model = Sequential()
    model.add(Dense(n_inputs, input_shape=(n_inputs, ), activation='relu'))
    model.add(Dense(32, activation='relu'))
    model.add(Dense(y_out, activation='softmax'))
    # Compile model
    model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
    return model
```

In [33]:

```
# Fitting and cross-validating ANN
estimator = KerasClassifier(build_fn=baseline_model, epochs=100, batch_size=5, verbose=0)
kfold = KFold(n_splits=5, shuffle=True)
results = cross_val_score(estimator, x, y, cv=kfold)
print("Neural Network Accuracy: %.2f%% (std. dev - %.2f%%)" % (results.mean()*100, results.std()*100))
```

Neural Network Accuracy: 74.86% (std. dev - 4.26%)